

Nikita Kazeev

kazeevn@gmail.com · [Website](#) · [GitHub](#) · [LinkedIn](#) · [Google Scholar](#) · [ORCID](#)

Research Scientist at the intersection of AI and Physics

WORK EXPERIENCE

CONSTRUCTOR GROUP *Machine Learning Consultant* 2021–present

NATIONAL UNIVERSITY OF SINGAPORE *Postdoc under [Kostya Novoselov](#)* 2022–present

Wyckoff Transformer, a generative model for materials [ICML 2025](#) · [GitHub](#)

- Developed a coordinate-free permutation-invariant autoregressive Transformer encoder model for material generation and property prediction based on symmetry inductive bias
- Increased materials discovery yield (S.U.N.) by 1.24x
- Accelerated inference by 4 orders of magnitude, to 0.05 GPU ms per structure
- Implemented the model in PyTorch; production-ready package published on [PyPI](#)
- Orchestrated large-scale (10k structures) DFT computations with VASP at an HPC cluster
- Led a team of 6

ML for defects in 2D crystals [npj 2D Mater. Appl.](#) · [npj Comput. Mater.](#) · [patent](#) · [GitHub](#)

- Conceived a novel sparse representation of crystals with defects
- Achieved a 3.7x energy prediction error reduction
- Optimised training and inference: 4x less memory, 8x fewer GPU operations
- Developed code for reproducible parallel running of machine learning experiments
- Led a team of 3

(CERN | YANDEX → HSE UNIVERSITY) *Researcher* 2014–2022

[2025 Breakthrough Prize in Fundamental Physics](#) as a member of LHCb.

Generative models uncertainty estimation [paper](#) · [conference](#) · [code 1](#) · [code 2](#)

- Developed the first methods for estimating uncertainty of conditional GANs.
- Led a team of 3 students.

Machine learning on noisy data [paper 1](#) · [paper 2](#) · [code](#)

- Developed a rigorous way to train ML algorithms on data with the label-noise model common in high-energy physics.

Muon identification at the LHCb experiment at CERN [paper](#) · [poster](#) · [IDAO-2019](#)

- Solved a classification problem over tabular data with noisy labels under timing constraints.
- Developed a gradient boosting model that reduced the false positive rate by 30% in the critical low-track-momentum region.
- Integrated the model into LHCb production, mostly in C++.
- Packaged the work as a data science competition problem for IDAO-2019.

CatBoost aka fighting biases with dynamic boosting [paper](#)

- Set up a distributed system for running experiments for the team with Bash and Python.
- Studied gradient boosting improvements with experiments on toy data.

EDUCATION

HSE University & Sapienza Università di Roma <i>PhD in Computer Science and Physics (Double Degree)</i> Supervisors: Andrey Ustyuzhanin & Barbara Sciascia Thesis: Machine Learning for particle identification in the LHCb detector	2016–2020
Product Management course at Yandex <i>Focused coursework in product strategy and execution.</i>	2018
Yandex School of Data Analysis <i>Master's-level CS programme</i> Algorithms, machine learning, deep learning, and distributed systems.	2013–2015
Moscow Institute of Physics and Technology <i>MS in Physics</i> Optimisation of data processing of the LHCb experiment.	2014–2016
<i>BS in Physics</i> Wave-packet molecular dynamics of electrons in nonideal plasma.	2010–2014
International Junior Science Olympiad (IJSO) Silver medal	Korea, 2008

MENTORSHIP

-
- Mentored 9 students ([1](#), [2](#), [3](#), [4](#), [5](#), [6](#), [7](#), [8](#), [9](#)), 3 interns, and student workshop projects ([1](#), [2](#)).

TEACHING & OUTREACH

-
- Taught at the Machine Learning in High Energy Physics Summer School in [2015](#), [2016](#), [2017](#), [2018](#), [2019](#), [2020](#), [2021](#), and [2023](#).
 - Lecturer in the [award-winning](#) Coursera Advanced Machine Learning [specialisation](#)
 - Taught and developed an introductory data analysis [course](#) at HSE.
 - Delivered more than 20 [conference talks](#).

SERVICE

-
- Reviewer for NeurIPS, RSC Advances, Machine Learning: Science and Technology, Digital Discovery
 - Peer Staff Supporter at NUS, serving as a first-line support for mental wellbeing.

SKILLS

- **Uncertainty quantification** — built the first methods for estimating the uncertainty of conditional GANs; calibrated generative and discriminative models.
- **ML on structured, non-text data** — tabular, graph, and atomistic representations, spanning gradient boosting to transformers.
- **Inference and systems at scale** — C++ production integration, large-scale DFT/VASP on HPC clusters, and reproducible experiment pipelines.
- **Physics domain depth** — high-energy and condensed-matter physics; first-principles intuition that shapes model design.

ML & AI • Python • PyTorch • C++ • Physics

SELECTED PUBLICATIONS

Nikita Kazeev, Wei Nong, Ignat Romanov, Ruiming Zhu, Andrey E Ustyuzhanin, Shuya Yamazaki, and Kedar Hipalgaoonkar. [“Wyckoff Transformer: Generation of Symmetric Crystals”](#). In: *Proceedings of the 42nd International Conference on Machine Learning*. Ed. by Aarti Singh, Maryam Fazel, Daniel Hsu, Simon Lacoste-Julien, Felix Berkenkamp, Tegan Maharaj, Kiri Wagstaff, and Jerry Zhu. Vol. 267. Proceedings of Machine Learning Research. PMLR, 13–19 Jul 2025, pp. 29495–29526.

Nikita Kazeev, A. R. Al-Maeeni, I. Romanov, et al. [“Sparse representation for machine learning the properties of defects in 2D materials”](#). In: *npj Computational Materials* 9 (2023), p. 113.

Maxim Borisyak and **Nikita Kazeev**. [“Machine Learning on data with sPlot background subtraction”](#). In: *Journal of Instrumentation* 14.08 (2019). Alphabetic order, I’m the corresponding author.

Denis Derkach, **Nikita Kazeev**, Fedor Ratnikov, Andrey Ustyuzhanin, and Alexandra Volokhova. [“Cherenkov detectors fast simulation using neural networks”](#). In: *Nuclear Instruments and Methods in Physics Research Section A* (2019). Alphabetic order, I’m the corresponding author.